

Learning Analytics- ML Model predictions

August 20, 2024

1 Learning Analytics and Model prediction

Goal is to predict the module completion rates using the courses, studentAssessment, and assessments datasets available on Open University Learning Analytics Dataset (OULAD) and investigate which model better predicts based on the model results.

Data Source: <https://www.kaggle.com/rocki37/open-university-learning-analytics-dataset>

Have used Python and SciKitLearn Machine Learning libraries for analyses of data from the OULAD Dataset, which contains data about courses, students and their interactions with Virtual Learning Environment (VLE) for seven selected courses (called modules).

Github repository# <https://github.com/shirleymassignments/Introduction-to-ML>

```
[77]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report
from scipy import stats
```

```
[64]: # Load datasets
courses = pd.read_csv('data/courses.csv')
student_assessments = pd.read_csv('data/studentAssessment.csv')
assessments = pd.read_csv('data/assessments.csv')
student_info = pd.read_csv('data/studentInfo.csv')

# Preview datasets
print("Courses data\n")
print(courses.head())
print("\nStudent_assessments data\n")
print(student_assessments.head())
print("\nAssessments data\n")
print(assessments.head())
```

```
print("\nStudent_info data\n")
print(student_info.head())
```

Courses data

	code_module	code_presentation	module_presentation_length
0	AAA	2013J	268
1	AAA	2014J	269
2	BBB	2013J	268
3	BBB	2014J	262
4	BBB	2013B	240

Student_assessments data

	id_assessment	id_student	date_submitted	is_banked	score
0	1752	11391	18	0	78.0
1	1752	28400	22	0	70.0
2	1752	31604	17	0	72.0
3	1752	32885	26	0	69.0
4	1752	38053	19	0	79.0

Assessments data

	code_module	code_presentation	id_assessment	assessment_type	date	weight
0	AAA	2013J	1752	TMA	19.0	10.0
1	AAA	2013J	1753	TMA	54.0	20.0
2	AAA	2013J	1754	TMA	117.0	20.0
3	AAA	2013J	1755	TMA	166.0	20.0
4	AAA	2013J	1756	TMA	215.0	30.0

Student_info data

	code_module	code_presentation	id_student	gender	region \
0	AAA	2013J	11391	M	East Anglian Region
1	AAA	2013J	28400	F	Scotland
2	AAA	2013J	30268	F	North Western Region
3	AAA	2013J	31604	F	South East Region
4	AAA	2013J	32885	F	West Midlands Region

	highest_education	imd_band	age_band	num_of_prev_attempts \
0	HE Qualification	90-100%	55<=	0
1	HE Qualification	20-30%	35-55	0
2	A Level or Equivalent	30-40%	35-55	0
3	A Level or Equivalent	50-60%	35-55	0
4	Lower Than A Level	50-60%	0-35	0

studied_credits disability final_result

0	240	N	Pass
1	60	N	Pass
2	60	Y	Withdrawn
3	60	N	Pass
4	60	N	Pass

```
[26]: ###Merge the datasets based on relevant keys to create a unified dataset for
      ↪analysis.
```

```
[68]: # Merge assessments with student assessments on assessment ID
merged_data = pd.merge(student_assessments, assessments, on='id_assessment')

# Merge the result with courses to get course information
merged_data = pd.merge(merged_data, courses, on=['code_module',
      ↪'code_presentation'])

merged_data = pd.merge(merged_data, student_info, on=['id_student',
      ↪'code_module', 'code_presentation'])
```

```
[69]: print(merged_data.describe())
```

	id_assessment	id_student	date_submitted	is_banked \
count	173912.000000	1.739120e+05	173912.000000	173912.000000
mean	26553.803556	7.051507e+05	116.032942	0.010977
std	8829.784254	5.523952e+05	71.484148	0.104194
min	1752.000000	6.516000e+03	-11.000000	0.000000
25%	15022.000000	5.044290e+05	51.000000	0.000000
50%	25359.000000	5.852080e+05	116.000000	0.000000
75%	34883.000000	6.344980e+05	173.000000	0.000000
max	37443.000000	2.698588e+06	608.000000	1.000000

	score	date	weight \
count	173739.000000	171047.000000	173912.000000
mean	75.799573	130.605623	12.743899
std	18.798107	78.025175	17.877301
min	0.000000	12.000000	0.000000
25%	65.000000	54.000000	0.000000
50%	80.000000	129.000000	9.000000
75%	90.000000	214.000000	18.000000
max	100.000000	261.000000	100.000000

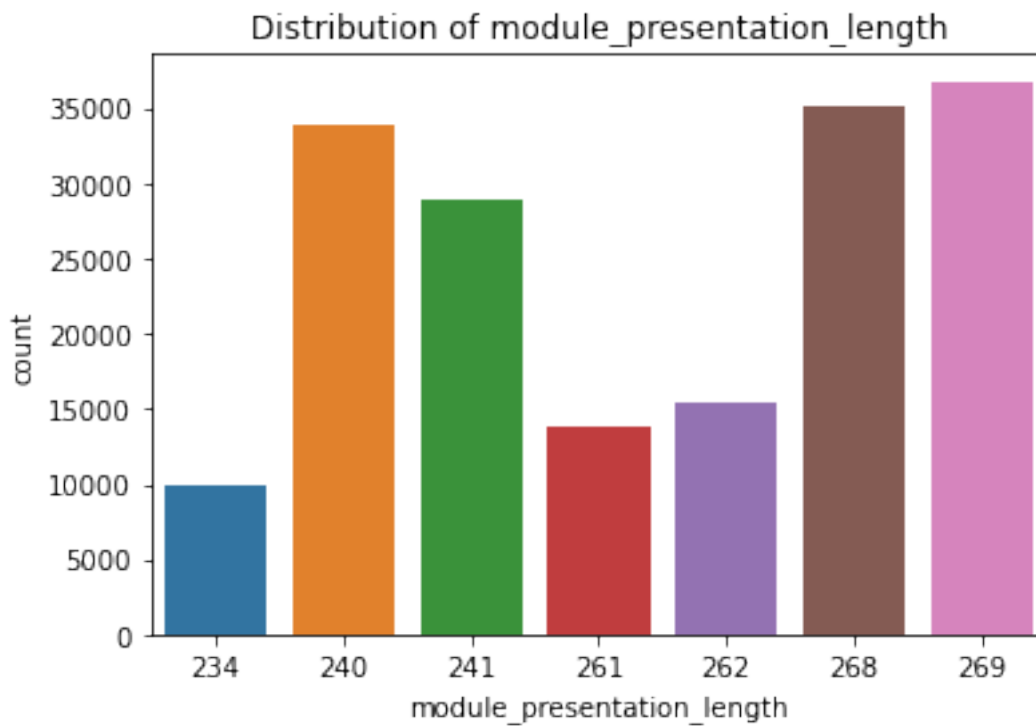
	module_presentation_length	num_of_prev_attempts	studied_credits
count	173912.000000	173912.000000	173912.000000
mean	255.233831	0.144907	76.779147
std	13.579401	0.445326	37.341741
min	234.000000	0.000000	30.000000
25%	240.000000	0.000000	60.000000
50%	262.000000	0.000000	60.000000

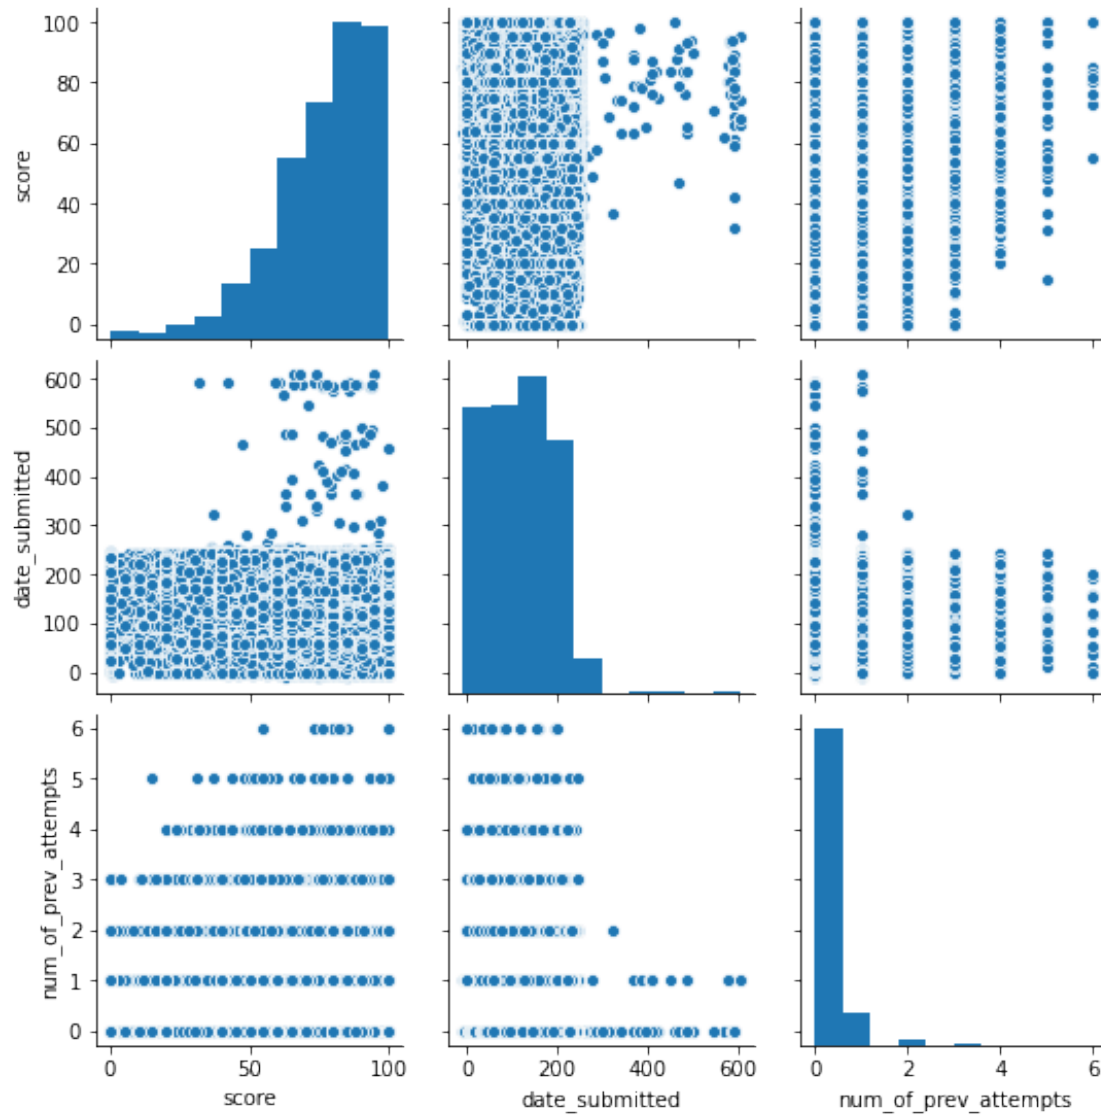
75%	268.000000	0.000000	90.000000
max	269.000000	6.000000	630.000000

```
[70]: ### Exploratory Data Analysis (EDA)
```

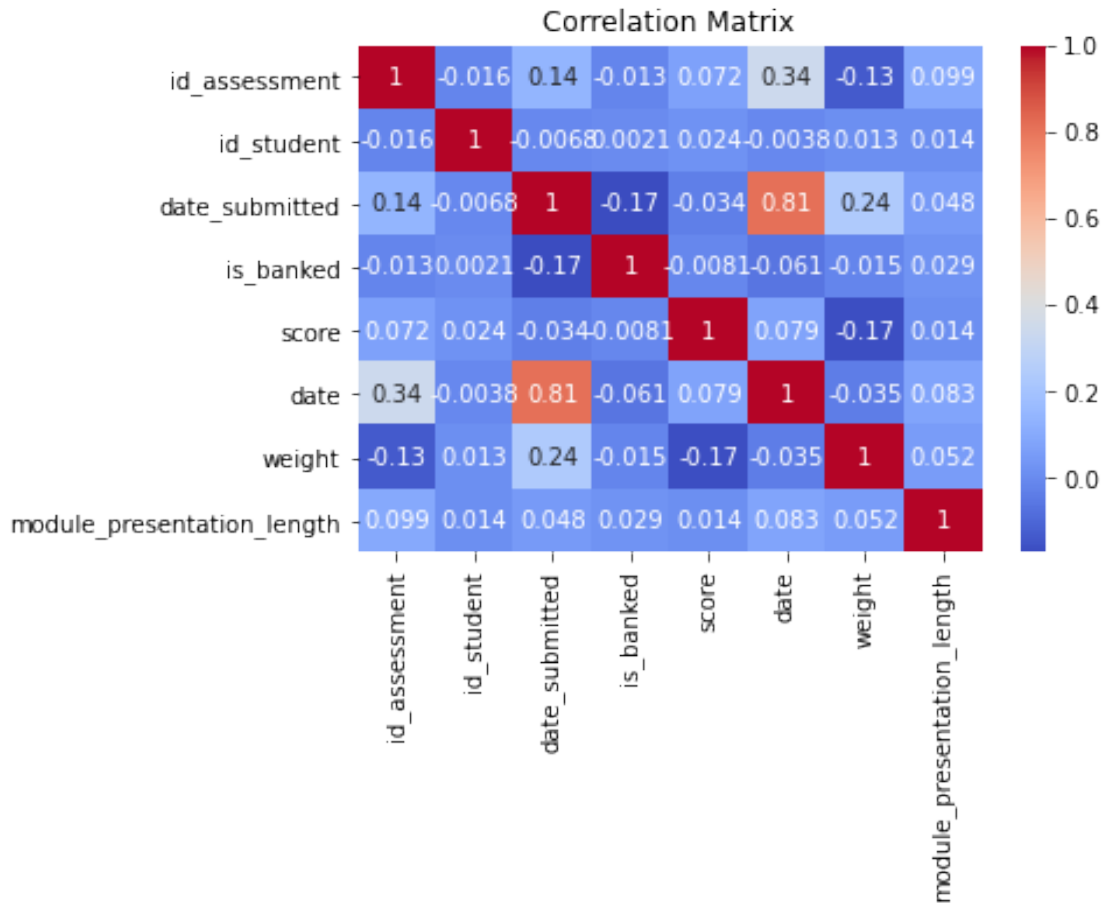
```
[74]: # Distribution of the target variable (e.g., 'final_result')
sns.countplot(merged_data['module_presentation_length'], data=merged_data)
plt.title('Distribution of module_presentation_length')
plt.show()

# Scatter plot for relationships
sns.pairplot(merged_data[['score', 'date_submitted', 'num_of_prev_attempts', '
↪ 'final_result']])
plt.show()
```





```
[52]: # Correlation matrix
corr_matrix = merged_data.corr()
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



The matrix suggests that factors like score, date_submitted, and weight have more substantial impacts on module completion rates

```
[53]: ###Data Cleaning
```

```
[75]: # Check for missing values
print(merged_data.isnull().sum())

# Handle missing values
merged_data = merged_data.dropna()

# Feature Engineering: Create aggregated features
# For example, calculate the total score per course per student
merged_data['total_score'] = merged_data.groupby(['id_student', 'code_module', 'code_presentation'])['score'].transform('sum')

# Drop duplicates or redundant columns if any
merged_data = merged_data.drop_duplicates(subset=['id_student', 'code_module', 'code_presentation'])
```

```

id_assessment          0
id_student              0
date_submitted         0
is_banked              0
score                  173
code_module             0
code_presentation      0
assessment_type        0
date                   2865
weight                 0
module_presentation_length 0
gender                 0
region                 0
highest_education      0
imd_band               7697
age_band               0
num_of_prev_attempts   0
studied_credits        0
disability             0
final_result           0
dtype: int64

```

```

[82]: # One-hot encoding for categorical features
categorical_features = ['code_module', 'code_presentation', 'gender', 'region',
↳ 'highest_education']
encoded_data = pd.get_dummies(merged_data[categorical_features],
↳ drop_first=True)

# Combine with the rest of the data
final_data = pd.concat([merged_data, encoded_data], axis=1)
final_data.drop(categorical_features, axis=1, inplace=True)

```

```

[56]: ## Data preparation

```

```

[83]: # Select predictors and target variable
X = final_data[['total_score', 'date', 'weight']] # Select relevant predictors
y = final_data['module_presentation_length'] # Continuous target: length of
↳ the module presentation

# Train-test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

# Standardization (if necessary)
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()

```

```
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

[57]: *## Model Building: Linear Regression*

```
[98]: # LinearRegression model
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test) # prediction

# Model Evaluation
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f'Mean Squared Error: {mse}')
print(f'R-squared: {r2}')

# Coefficients interpretation
coefficients = pd.DataFrame(model.coef_, X.columns, columns=['Coefficient'])
print(coefficients)

# Logistic Regression Model
log_reg = LogisticRegression(max_iter=1000)
log_reg.fit(X_train, y_train)
y_pred_lr = log_reg.predict(X_test)

# Model Evaluation
print("\n\nLogistic Regression:\n\n")
print(f'Accuracy: {accuracy_score(y_test, y_pred_lr)}')
print(confusion_matrix(y_test, y_pred_lr))
print(classification_report(y_test, y_pred_lr))
```

Mean Squared Error: 168.54732316882414

R-squared: 0.020088379419650915

	Coefficient
total_score	-1.394672
date	0.808240
weight	1.014224

Logistic Regression:

Accuracy: 0.4300100704934542

```
[[244  0  0  0  0  0  0]
 [  0 278 52  0 77 252 55]
 [  0  0 156 29 79 161 434]]
```



```

[ 0  0 176 32  0  0 240]
[ 0 103  3  0 504  0  3]
[ 0 292  1  2  93 431 181]
[ 0  0 154 29 99 315 490]]
      precision    recall  f1-score   support

    234         1.00      1.00      1.00        244
    240         0.41      0.39      0.40        714
    241         0.29      0.18      0.22        859
    261         0.35      0.07      0.12        448
    262         0.59      0.82      0.69        613
    268         0.37      0.43      0.40       1000
    269         0.35      0.45      0.39       1087

 accuracy                   0.43       4965
 macro avg              0.48      0.48      0.46       4965
 weighted avg          0.41      0.43      0.41       4965

```

```

[99]: # Random Forest Model
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
y_pred_rf = rf_model.predict(X_test)

print("\n\nRandom Forest:\n\n")
print(f'Accuracy: {accuracy_score(y_test, y_pred_rf)}')
print(confusion_matrix(y_test, y_pred_rf))
print(classification_report(y_test, y_pred_rf))

```

Random Forest:

```

Accuracy: 0.611681772406848
[[244  0  0  0  0  0  0]
 [  0 408  0  1  0 304  1]
 [  0  1 306 167  0  33 352]
 [  0  0 135 278  0  0  35]
 [  0  0  0  0 612  0  1]
 [  0 289 27  0  0 563 121]
 [  0  0 289 69  0 103 626]]
      precision    recall  f1-score   support

    234         1.00      1.00      1.00        244
    240         0.58      0.57      0.58        714
    241         0.40      0.36      0.38        859

```

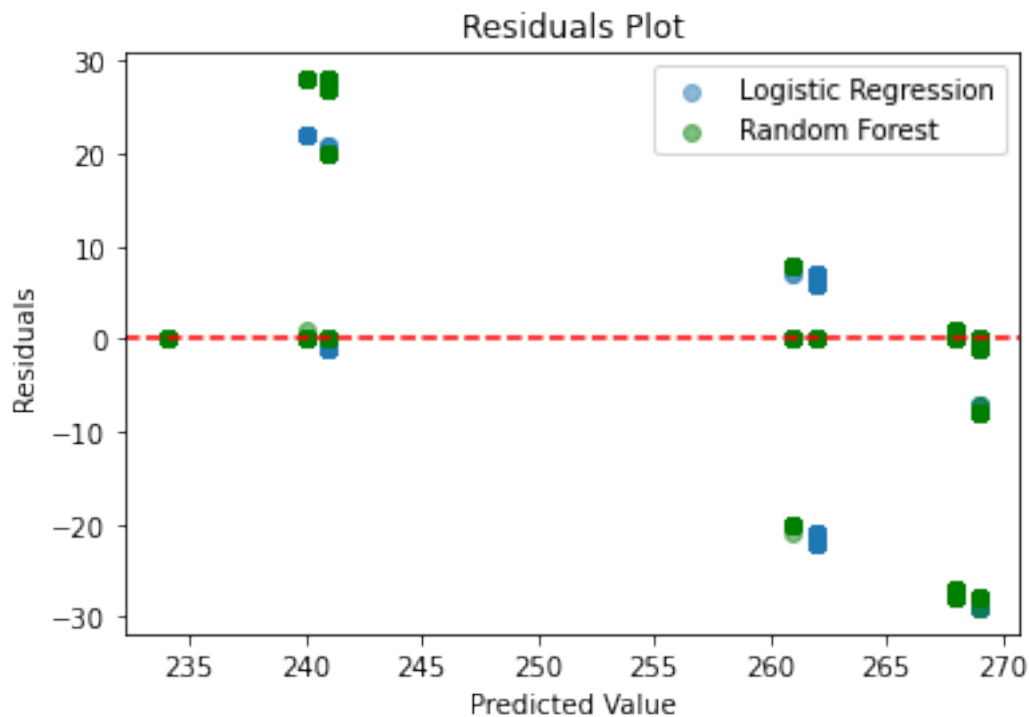
261	0.54	0.62	0.58	448
262	1.00	1.00	1.00	613
268	0.56	0.56	0.56	1000
269	0.55	0.58	0.56	1087
accuracy			0.61	4965
macro avg	0.66	0.67	0.67	4965
weighted avg	0.61	0.61	0.61	4965

```
[97]: ### Residual Analysis
```

```
[96]: # Residuals for Logistic Regression
residuals_lr = y_test - y_pred_lr

# Residuals for Random Forest
residuals_rf = y_test - y_pred_rf

# Plot Residuals
plt.figure()
plt.scatter(y_pred_lr, residuals_lr, alpha=0.5, label='Logistic Regression')
plt.scatter(y_pred_rf, residuals_rf, alpha=0.5, color='green', label='Random_
↳Forest')
plt.axhline(0, color='red', linestyle='--')
plt.xlabel('Predicted Value')
plt.ylabel('Residuals')
plt.title('Residuals Plot')
plt.legend()
plt.show()
```



2 Analysis

3 Linear Regression

MSE indicates the average squared difference between predicted and actual values. A higher MSE suggests that the model's predictions are not close to the actual values. R-squared is very low (0.020), meaning the model explains only 2% of the variance in the target variable. This suggests that the linear regression model is not suitable for this data.

4 Logistic Regression

Accuracy is 0.43, which is quite low and suggests that the model performs poorly in predicting module completion.

Precision and Recall vary significantly across classes, with high precision and recall for some classes (ex: Class 234) and very low for others (Ex: Class 261)

5 Random Forest


Accuracy is 0.61, indicating that the Random Forest model performs better than Logistic Regression. Precision and Recall are more balanced across classes compared to Logistic Regression.


6 Visualizing Model Performance

```
[101]: from sklearn.metrics import ConfusionMatrixDisplay

# Logistic Regression Confusion Matrix
cm_lr = confusion_matrix(y_test, y_pred_lr)
disp_lr = ConfusionMatrixDisplay(confusion_matrix=cm_lr, display_labels=['234', '240', '241', '261', '262', '268', '269'])
disp_lr.plot(cmap=plt.cm.Blues)
plt.title('Logistic Regression Confusion Matrix')
plt.show()

# Random Forest Confusion Matrix
cm_rf = confusion_matrix(y_test, y_pred_rf)
disp_rf = ConfusionMatrixDisplay(confusion_matrix=cm_rf, display_labels=['234', '240', '241', '261', '262', '268', '269'])
disp_rf.plot(cmap=plt.cm.Blues)
plt.title('Random Forest Confusion Matrix')
plt.show()
```





7 Summary and Conclusion

Based on the model results, it concludes that the Logistic Regression model shows better performance than Linear Regression with an accuracy of 0.43. However, it struggles with class imbalance, and its precision/recall are inconsistent across classes. The Random Forest model has the highest accuracy (0.61) and provides more balanced precision and recall across different classes compared to Logistic Regression.

Random Forest is the best-performing model among the three for predicting module completion rates. It achieves the highest accuracy and shows more balanced performance across different classes.